

RAG PDF Summarization Test Document

This is a sample PDF created to test the RAG summarization capabilities.

Retrieval-Augmented Generation (RAG) is a technique for enhancing the accuracy and reliability of large language models by combining them with external knowledge sources.

In this example, we are using LangChain, ChromaDB, and Ollama.

LangChain provides the framework, ChromaDB acts as the vector store, and Ollama runs the LLM.

The goal is to summarize this content effectively.